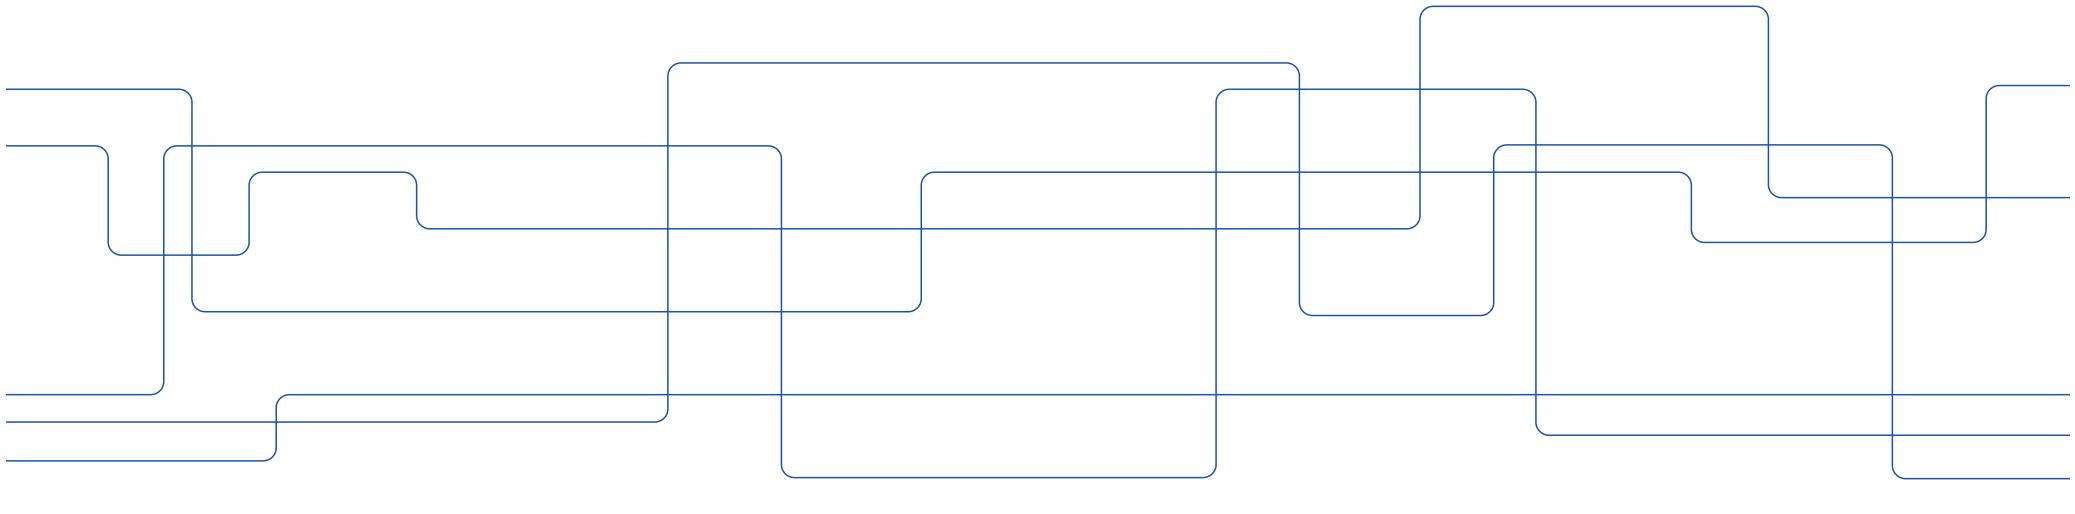




Chapter 6: The Link Layer and LANs

IK1203

Peter Sjödin, psj@kth.se





Chapter 6

The Link Layer and LANS

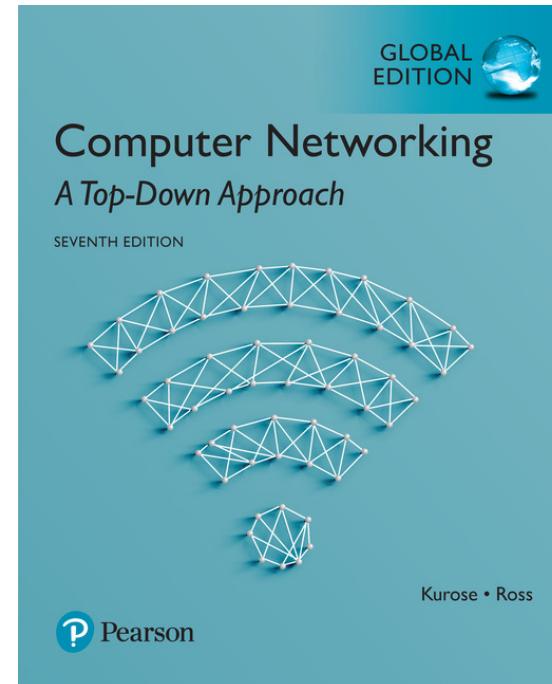
A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a lot of work on our part. In return for use, we only ask the following:

- ❖ If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- ❖ If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2012
J.F Kurose and K.W. Ross, All Rights Reserved



*Computer
Networking: A Top
Down Approach*

7th Edition, Global Edition
Jim Kurose, Keith Ross
Pearson
April 2016



Chapter 6: Link layer

our goals:

- ❖ understand principles behind link layer services:
 - error detection and correction
 - sharing a broadcast channel: multiple access
 - link layer addressing
 - local area networks: Ethernet, VLANs
- ❖ Instantiation and implementation of various link layer technologies



Link layer, LANs: outline

6.1 introduction, services

**6.2 error detection,
correction**

**6.3 multiple access
protocols**

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

**6.5 link virtualization:
MPLS**

**6.6 data center
networking**

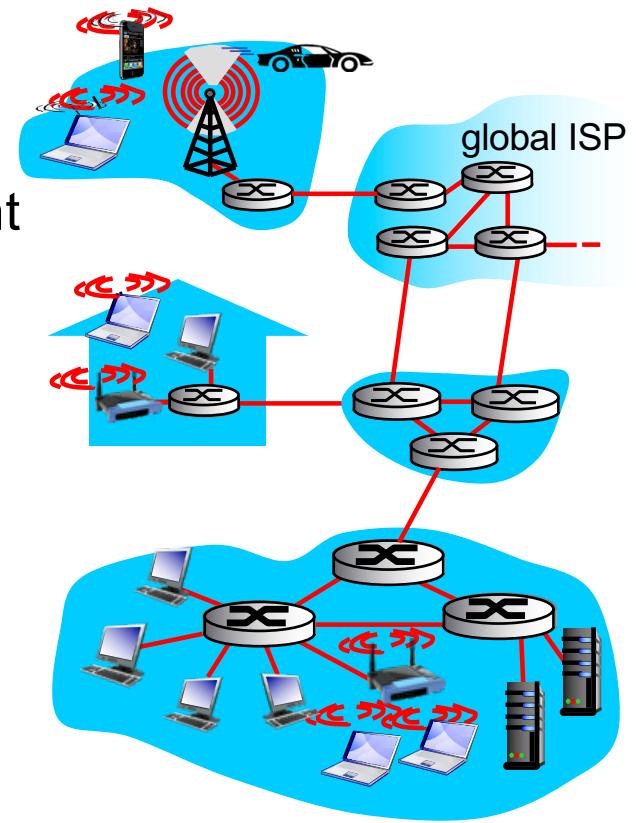
**6.7 a day in the life of a
web request**

Link layer: introduction

terminology:

- ❖ hosts and routers: **nodes**
- ❖ communication channels that connect adjacent nodes along communication path: **links**
 - wired links
 - wireless links
 - LANs
- ❖ layer-2 packet: **frame**, encapsulates datagram

link layer has the responsibility of transferring datagram from one node to *physically adjacent* node over a link





Link layer: context

- ❖ datagram transferred by different link protocols over different links:
 - e.g., Ethernet on first link, frame relay on intermediate links, 802.11 on last link
- ❖ each link protocol provides different services
 - e.g., may or may not provide reliable data transfer over link

transportation analogy:

- ❖ trip from Princeton to Lausanne
 - limo: Princeton to JFK
 - plane: JFK to Geneva
 - train: Geneva to Lausanne
- ❖ tourist = **datagram**
- ❖ transport segment = **communication link**
- ❖ transportation mode = **link layer protocol**
- ❖ travel agent = **routing algorithm**



Link layer services

- *framing, link access:*
 - encapsulate datagram into frame, adding header and trailer
 - channel access, if shared medium
 - “MAC” addresses used in frame headers to identify source and destination
 - > “Medium Access Control”
 - > *different from IP addresses!*
- *reliable delivery between adjacent nodes*
 - we learned how to do this already (chapter 3)!
 - seldom used on low bit-error link (fiber, some twisted pair)
 - wireless links: high error rates
 - > *Q: why both link-level and end-end reliability?*



Link layer services (more)

❖ *flow control:*

- pacing between adjacent sending and receiving nodes

❖ *error detection:*

- errors caused by signal attenuation and noise
- receiver detects presence of errors:
 - > *signals sender for retransmission or drops frame*

❖ *error correction:*

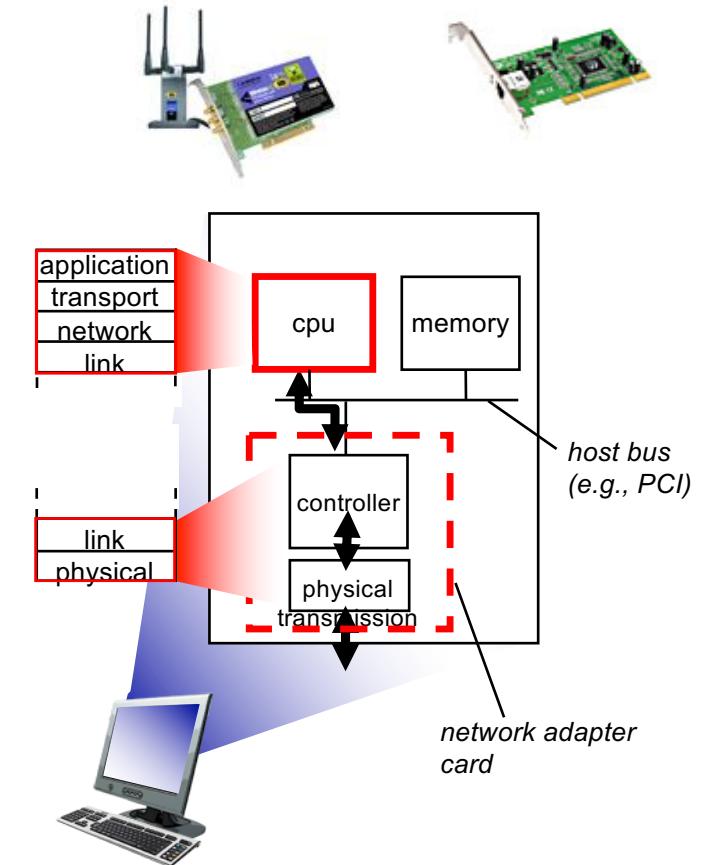
- receiver identifies *and corrects* bit error(s) without resorting to retransmission

❖ *half-duplex and full-duplex*

- with half duplex, nodes at both ends of link can transmit, but not at same time

Where is the link layer implemented?

- in each and every host
- link layer implemented in “adaptor” (aka *network interface card* or NIC) or on a chip
 - Ethernet card, 802.11 card; Ethernet chipset
 - implements link and physical layer
- attaches into host’s system buses
- combination of hardware, software, and firmware





Link layer, LANs: outline

6.1 introduction, services

6.2 error detection,
correction

6.3 multiple access
protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization:
MPLS

6.6 data center
networking

6.7 a day in the life of a
web request

Error detection

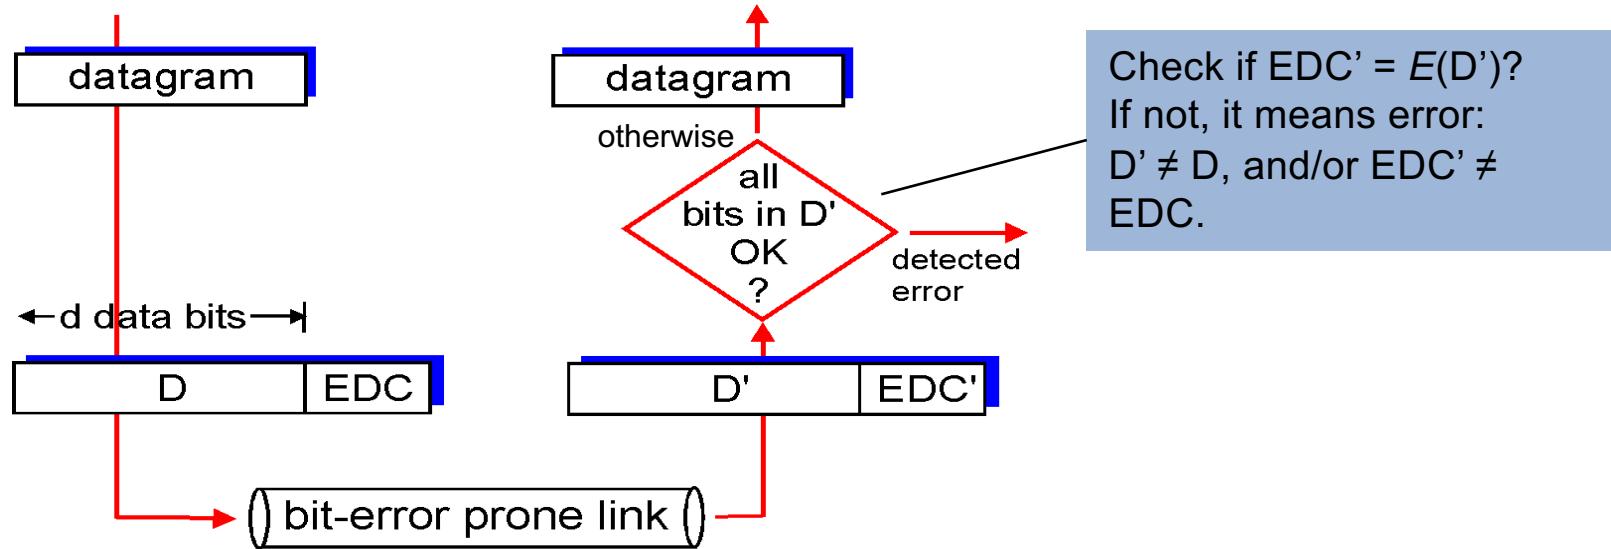
- Receiver cannot tell by just looking at the data whether data has been corrupted
 - Need to transfer more information to help with this – redundancy!

D: Data protected by error detection

$E()$: Error detection function

EDC: Error Detection and Correction bits, sent together with data, calculated as:

$$\text{EDC} = E(D)$$





Error detection functions

- Error detection not 100% reliable!
 - protocol may miss some errors, but rarely
 - larger EDC field yields better error detection and correction

Name	EDC Length	Description	Examples
Even parity	1 bit	Set so total number of one's is even. Detects single bit errors.	Memory, RAID disks
Internet checksum	16 bits	16-bit sum (modular arithmetic)	IP, TCP, UDP
CRC-32	32 bits	Cyclic code (polynomial division). Detects all burst errors up to 32 bits long (and some longer).	Ethernet



Link layer, LANs: outline

6.1 introduction, services

6.2 error detection,
correction

6.3 multiple access
protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization:
MPLS

6.6 data center
networking

6.7 a day in the life of a
web request

Multiple access links, protocols

two types of “links”:

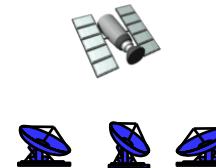
- point-to-point
 - PPP for dial-up access
 - point-to-point link between Ethernet switch and host
- *broadcast (shared wire or medium)*
 - old-fashioned Ethernet
 - upstream HFC (Hybrid Fiber-Coaxial cable – cable Internet)
 - 802.11 wireless LAN



shared wire (e.g.,
cabled Ethernet)



shared RF
(e.g., 802.11 WiFi)



shared RF
(satellite)



humans at a
cocktail party
(shared air, acoustical)



Multiple access protocols

- ❖ single shared broadcast channel
- ❖ two or more simultaneous transmissions by nodes: interference
 - *collision* if a node receives two or more signals at the same time

multiple access protocol

- ❖ distributed algorithm that determines how nodes share a channel, i.e., determines when a node can transmit



An ideal multiple access protocol

given: broadcast channel of rate R bit/s

desiderata:

1. when one node wants to transmit, it can send at rate R
2. when M nodes want to transmit, each can send at average rate R/M
3. fully decentralized:
 - > *no special node to coordinate transmissions*
 - > *no synchronization of clocks*
4. simple



MAC protocols: taxonomy

three broad classes:

- *channel partitioning*
 - divide channel into smaller "pieces" (time slots, frequency, code)
 - allocate piece to node for exclusive use
 - TDMA (Time Division Multiple Access), FDMA, CDMA
- *random access*
 - channel not divided
 - > *collisions could happen*
 - > *"recover" from collisions*
 - Ethernet, IEEE 802.11
- *"taking turns"*
 - nodes take turns
 - > *but nodes with more to send can take longer turns*
 - coordinated access
 - Bluetooth, token ring, IEEE 802.11 PCF (Point Coordination Function)



Random access protocols

- ❖ when node has packet to send
 - transmit at full channel data rate R
 - no *a priori* coordination among nodes
- ❖ two or more transmitting nodes → collision
- ❖ random access MAC protocol specifies:
 - how to detect collisions
 - how to recover from collisions (e.g., via delayed retransmissions)
- ❖ examples of random access MAC protocols:
 - slotted ALOHA
 - ALOHA
 - CSMA, CSMA/CD, CSMA/CA



CSMA (carrier sense multiple access)

CSMA: sense (listen) before transmit:

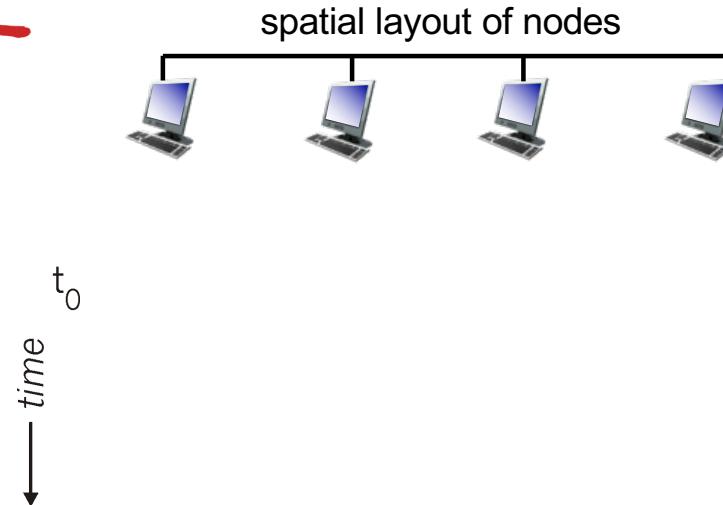
if channel sensed idle: transmit entire frame

- if channel sensed busy, defer transmission
- human analogy: don't interrupt others!



CSMA collisions

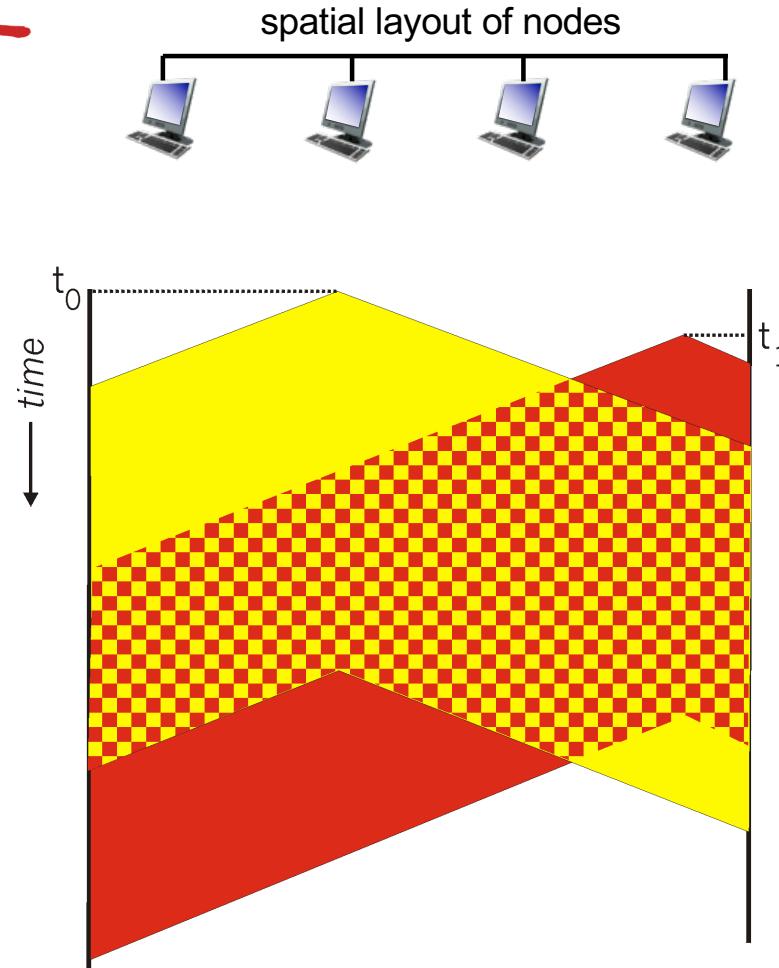
- **collisions can still occur:**
propagation delay means two nodes may not hear each other's transmission
- **collision:** entire packet transmission time wasted
 - distance and propagation delay play role in determining collision probability





CSMA collisions

- **collisions can still occur:**
propagation delay means two nodes may not hear each other's transmission
- **collision:** entire packet transmission time wasted
 - distance and propagation delay play role in determining collision probability



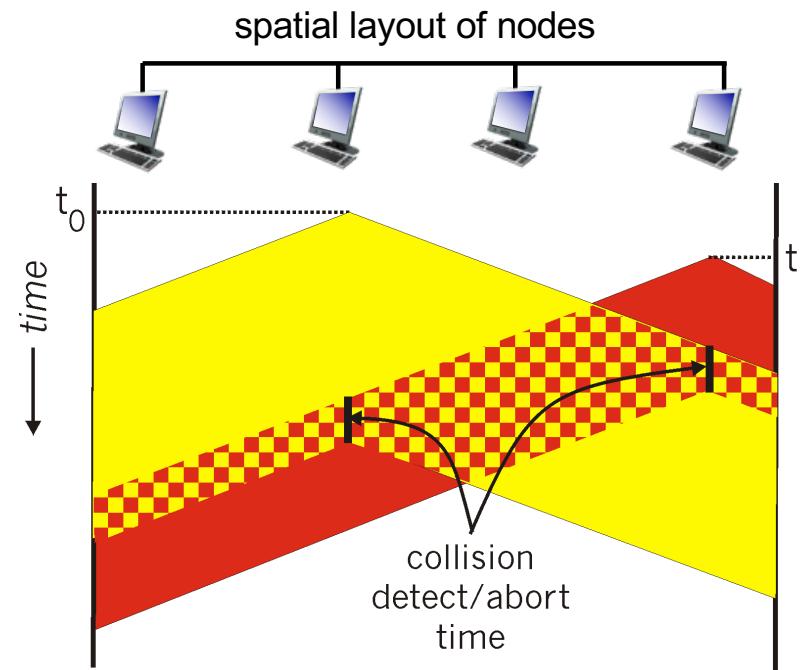


CSMA/CD (collision detection)

CSMA/CD: carrier sensing, deferral as in CSMA

- collisions *detected* within short time
 - colliding transmissions aborted, reducing channel wastage
- ❖ collision detection:
- easy in wired LANs: measure signal strengths, compare transmitted and received signals
 - difficult in wireless LANs: received signal strength overwhelmed by local transmission strength
- ❖ human analogy: the polite conversationalist

CSMA/CD (collision detection)





Summary of MAC protocols

- ❖ *channel partitioning*, by time, frequency or code
 - Time Division, Frequency Division
- ❖ *random access* (dynamic),
 - ALOHA, S-ALOHA, CSMA, CSMA/CD
 - carrier sensing: easy in some technologies (wire), hard in others (wireless)
 - CSMA/CD used in Ethernet
 - CSMA/CA used in 802.11
- ❖ *taking turns*
 - polling from central site, token passing
 - bluetooth, FDDI, 802.11 with PCF (Point Coordination Function), token ring



Link layer, LANs: outline

- 6.1 introduction, services
- 6.2 error detection, correction
- 6.3 multiple access protocols
- 6.4 LANs
 - addressing, ARP
 - Ethernet
 - switches
 - VLANs
- 6.5 link virtualization: MPLS
- 6.6 data center networking
- 6.7 a day in the life of a web request

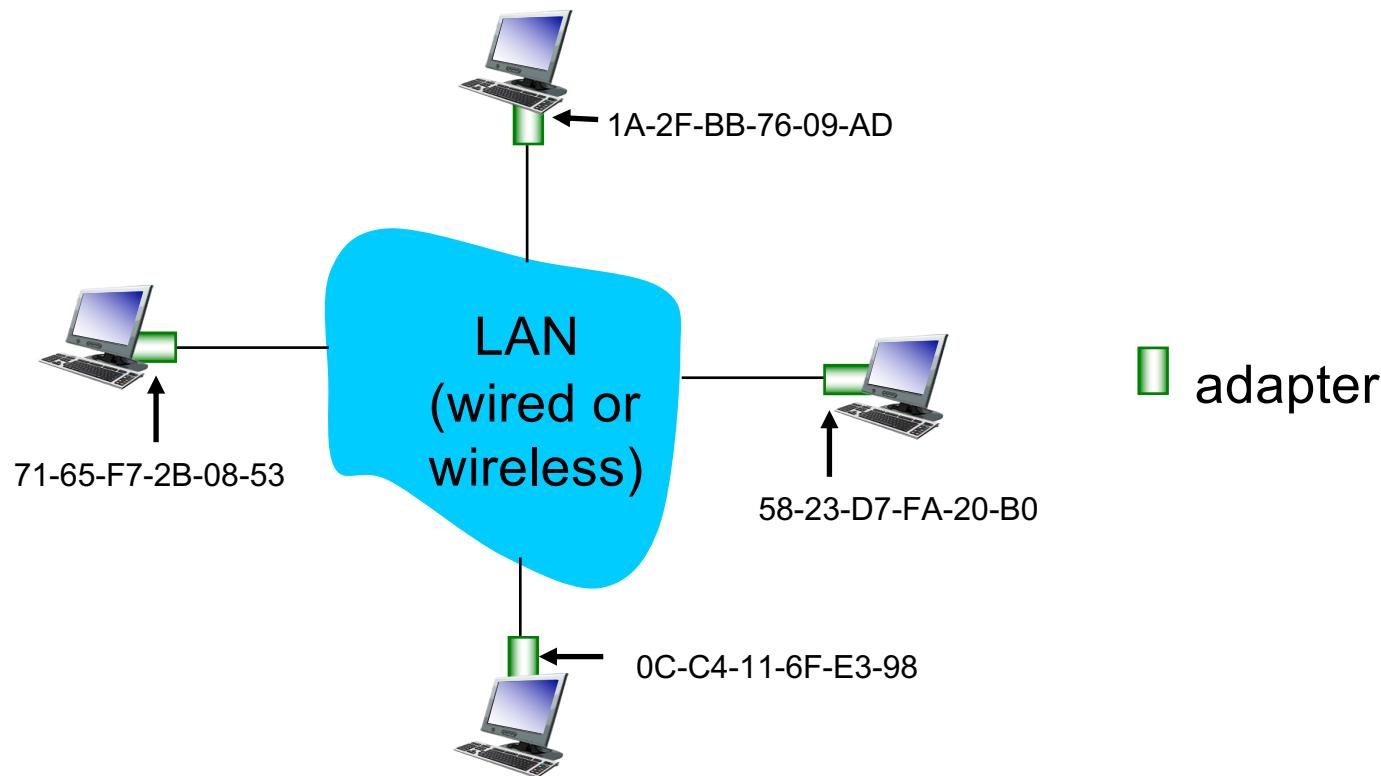


MAC addresses and ARP

- 32-bit (128-bit) IP address:
 - *network-layer* address for interface
 - used for layer 3 (network layer) forwarding
- MAC (or LAN or physical or Ethernet) address:
 - function: *used “locally” to get frame from one interface to another physically-connected interface (same network, in IP-addressing sense)*
 - 48-bit MAC address (for most LANs) burned in NIC ROM, also sometimes software configurable
 - e.g.: 0A-2F-BB-76-09-AD (often also written as 0A:2F:BB:76:09:AD)
 - /
 - hexadecimal (base 16) notation
 - (each “number” (letter/digit) represents 4 bits)

LAN addresses and ARP

each adapter on LAN has unique *LAN* address



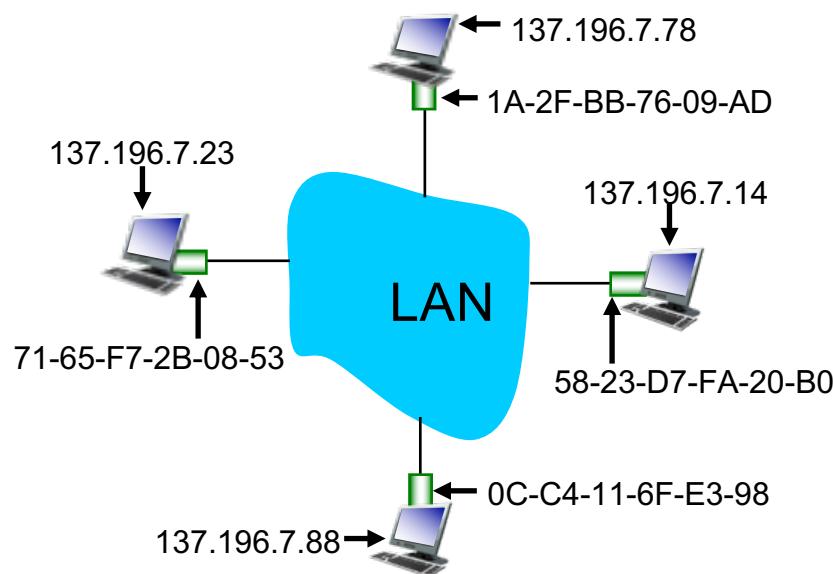


LAN addresses (more)

- MAC address allocation administered by IEEE
 - manufacturer buys portion of MAC address space
- MAC address bound to specific hardware (adapter)
- analogy:
 - MAC address: like Social Security Number
 - > “*personnummer*”
 - IP address: like postal address
- MAC flat address → portability
 - can move LAN card from one LAN to another
- IP hierarchical address *not* portable
 - address depends on IP subnet to which node is attached

ARP: address resolution protocol

Question: how to determine an interface's MAC address, knowing its IP address?



ARP table: each IP node (host or router) on a LAN has a table:

- IP/MAC address mappings for nodes on the LAN:
< IP address; MAC address; TTL >
- > *TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)*



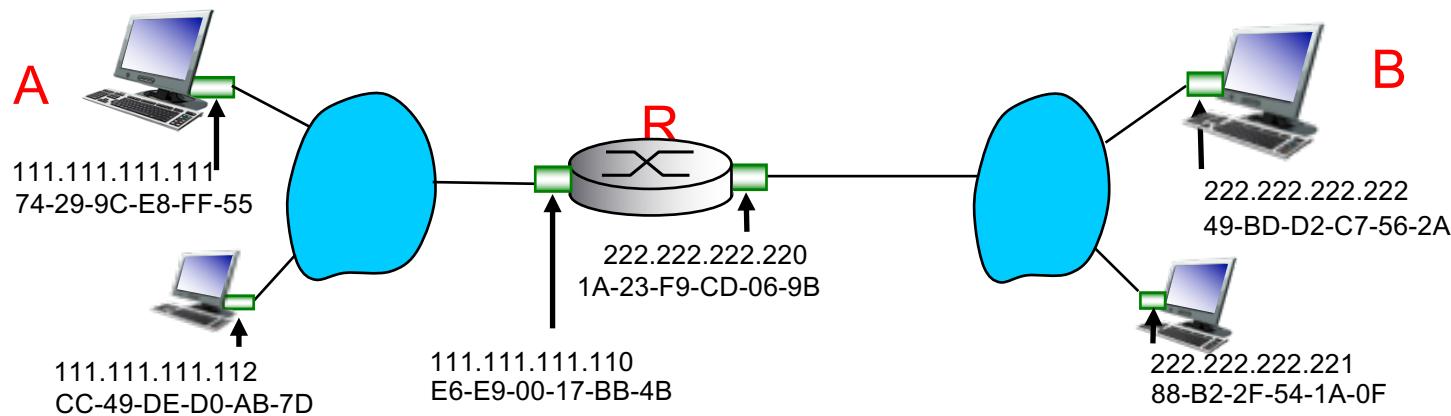
ARP protocol: same LAN

- A wants to send datagram to B
 - B's MAC address not in A's ARP table.
- A **broadcasts** ARP query, containing B's IP address
 - To broadcast MAC address = FF-FF-FF-FF-FF-FF
 - all nodes on LAN receive the ARP query
- B receives ARP query, replies to A with its (B's) MAC address
 - frame sent to A's MAC address (unicast)
- A caches (saves) IP-to-MAC address pair in its ARP table until information becomes old (times out)
 - soft state: information that times out (goes away) unless refreshed
- ARP is "plug-and-play":
 - nodes create their ARP tables *without configuration or intervention from network administrator*

Addressing: routing to another LAN

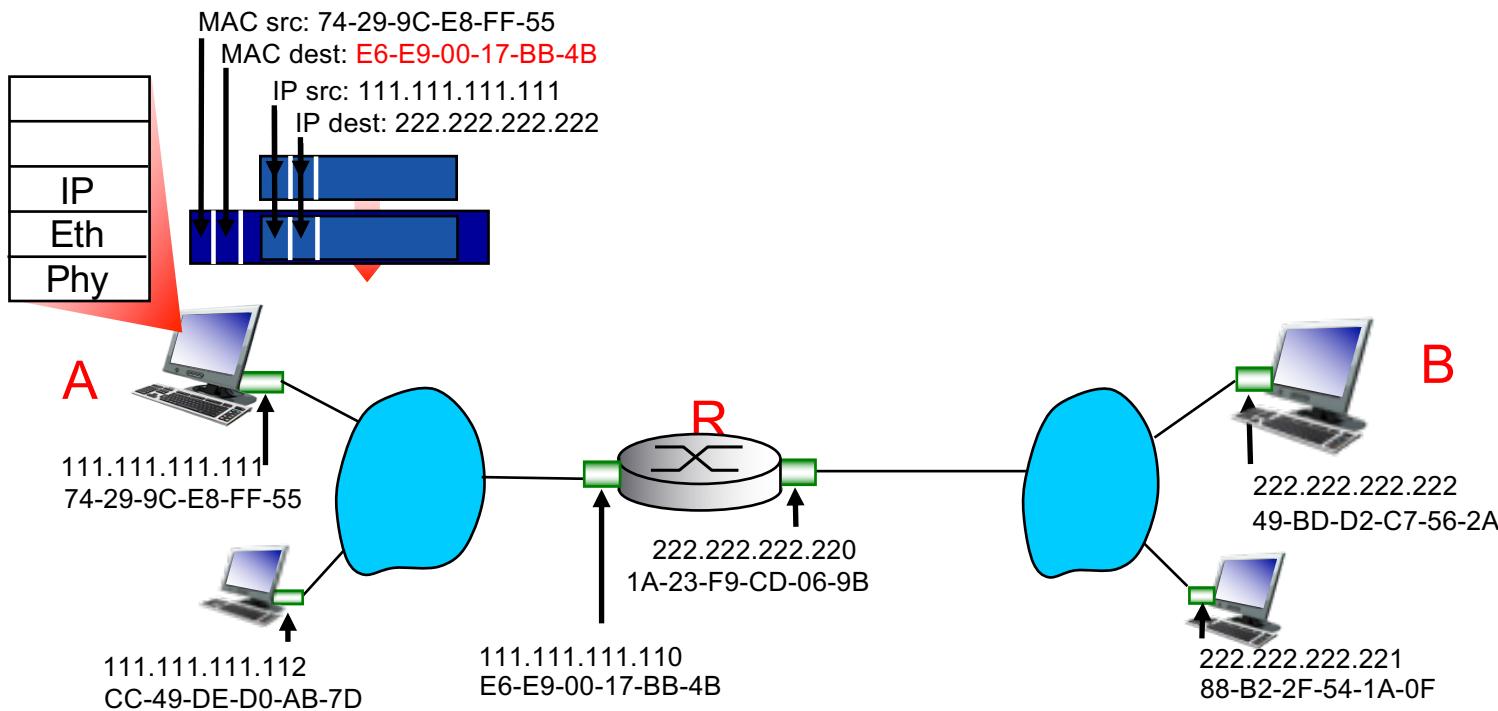
walkthrough: **send datagram from A to B via R**

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R (how?)
- assume A knows R's MAC address (how?)



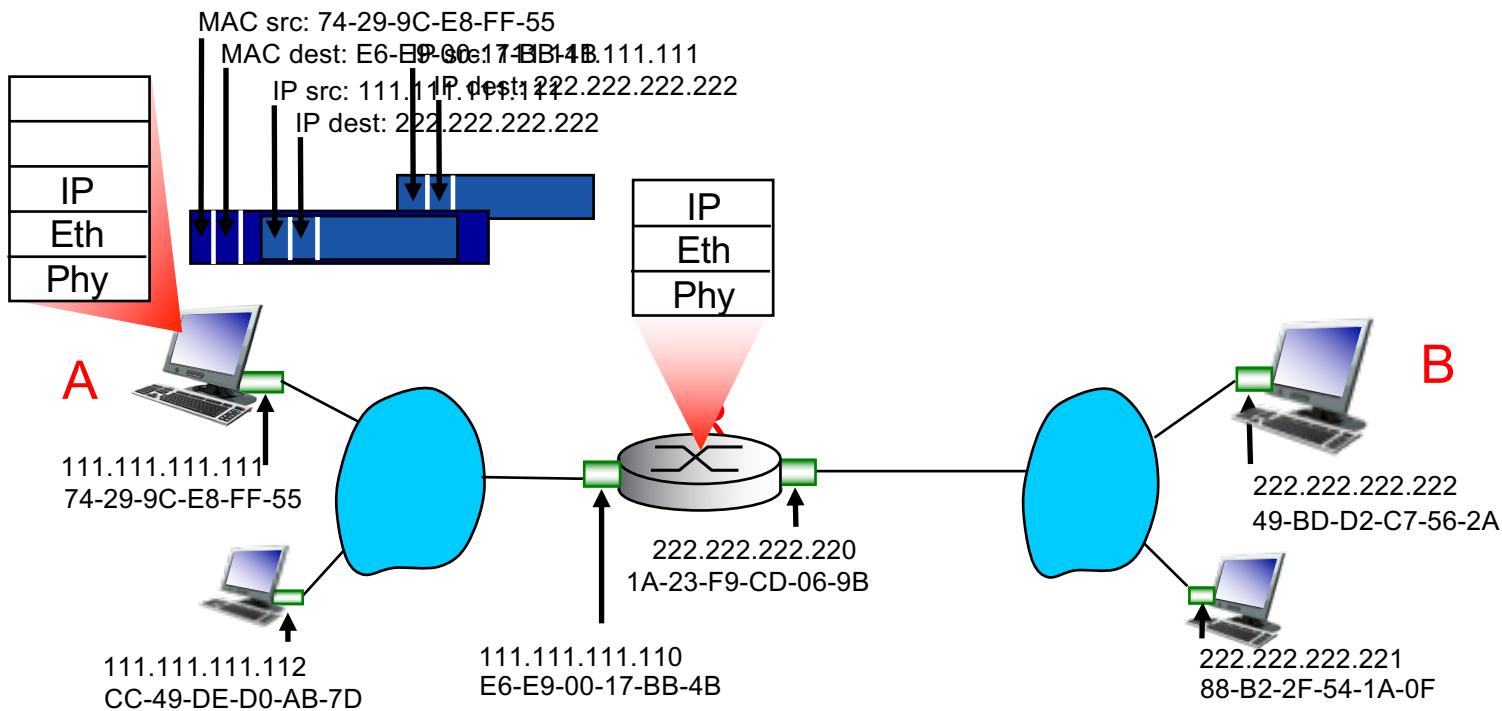
Addressing: routing to another LAN

- ❖ A creates “A-to-B” IP datagram, with IP source A and destination B
- ❖ A creates link-layer frame with R's MAC address as destination, frame contains “A-to-B” IP datagram



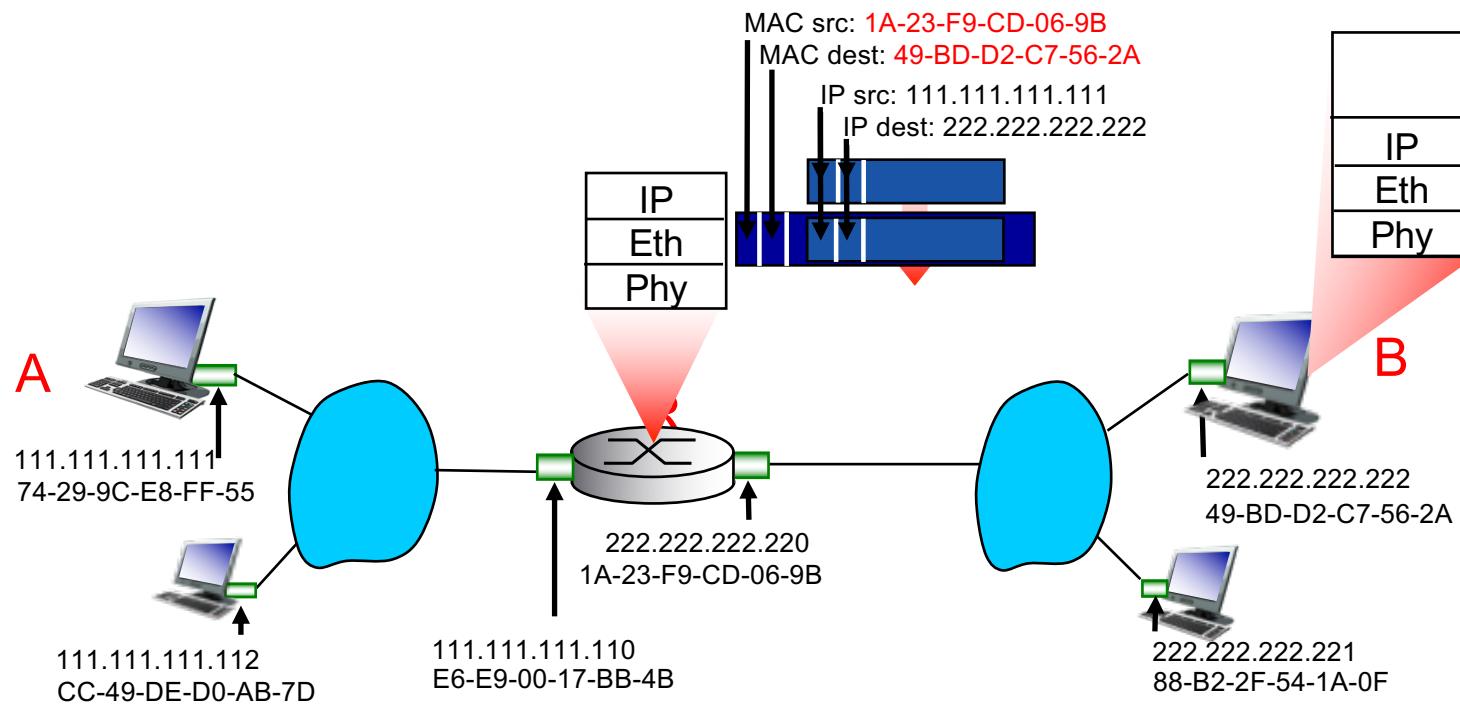
Addressing: routing to another LAN

- ❖ frame sent from A to R
- ❖ frame received at R, datagram removed, passed up to IP



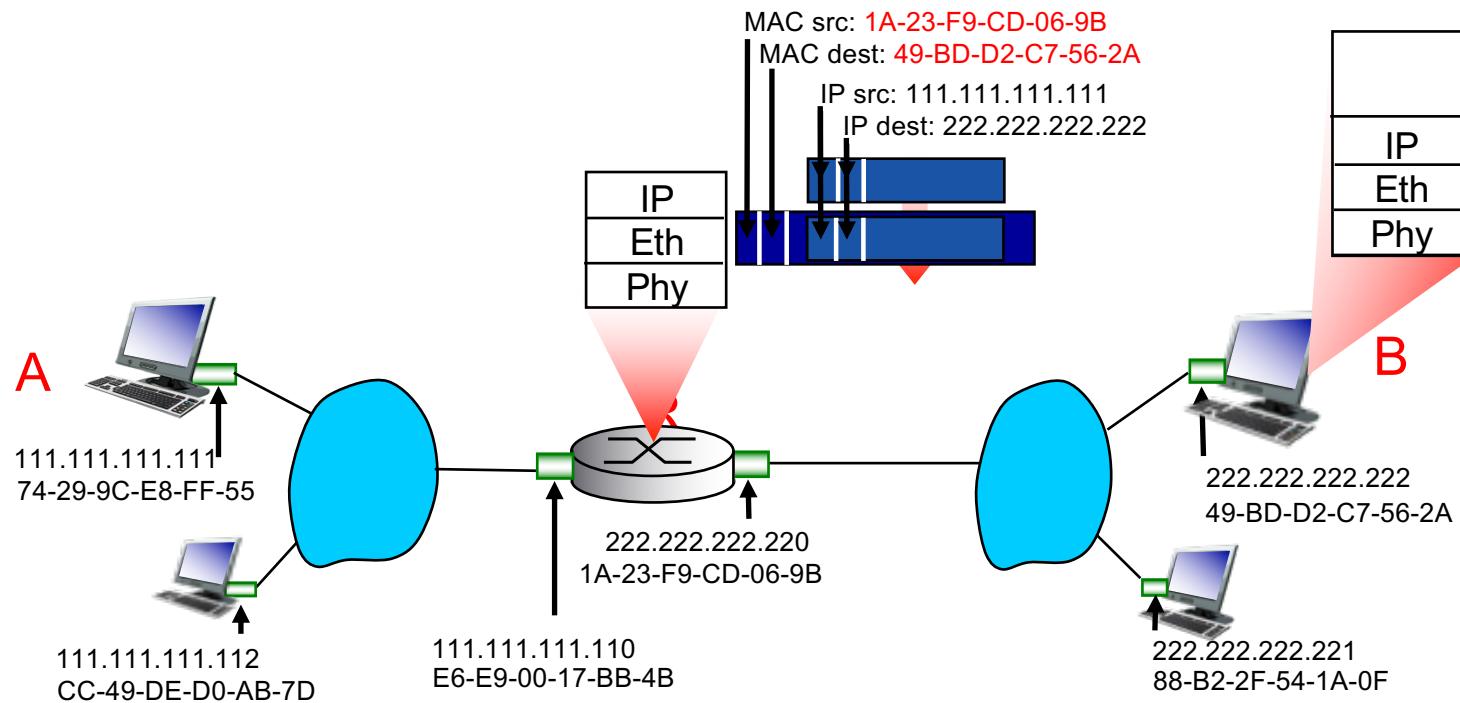
Addressing: routing to another LAN

- ❖ R forwards “A-to-B” datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as destination, frame contains “A-to-B” IP datagram



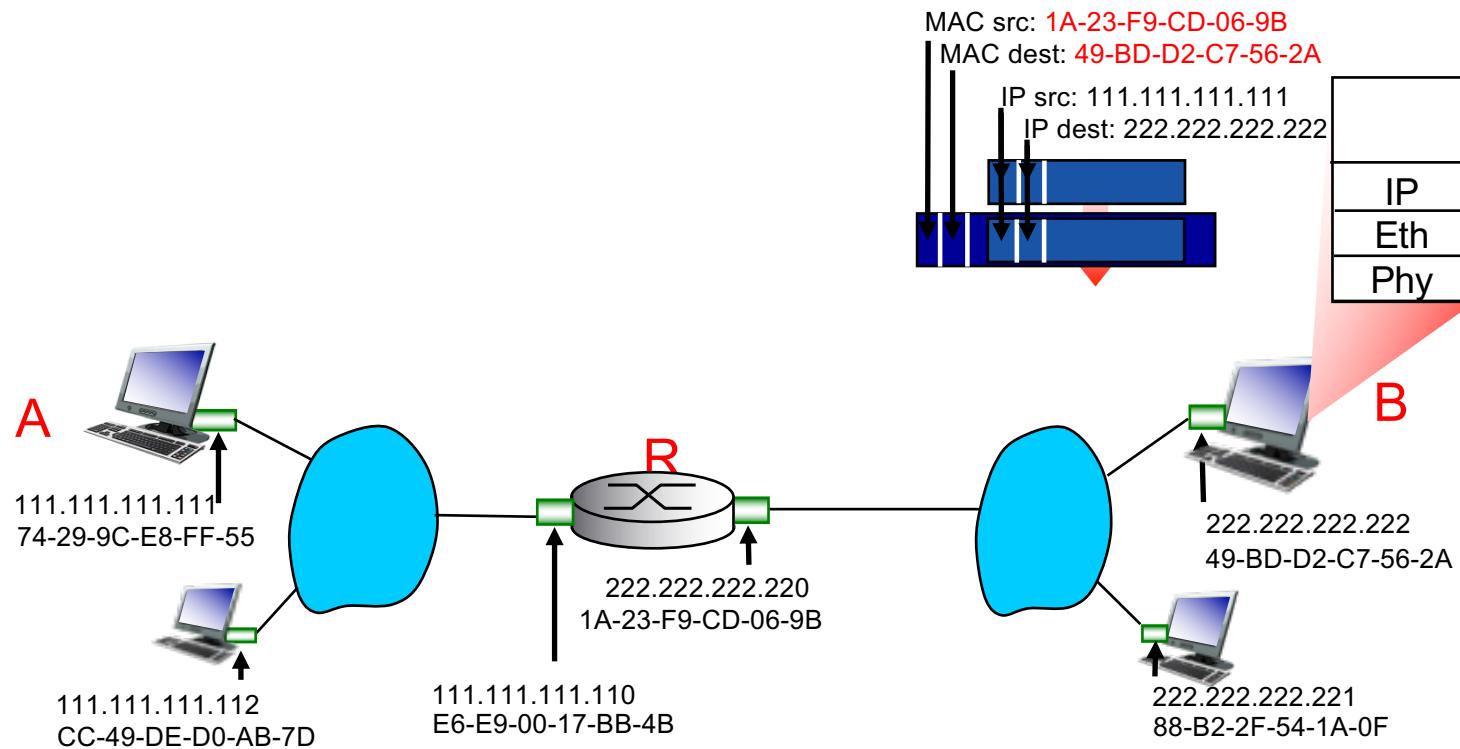
Addressing: routing to another LAN

- ❖ R forwards “A-to-B” datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as destination, frame contains “A-to-B” IP datagram



Addressing: routing to another LAN

- ❖ R forwards “A-to-B” datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as destination, frame contains “A-to-B” IP datagram





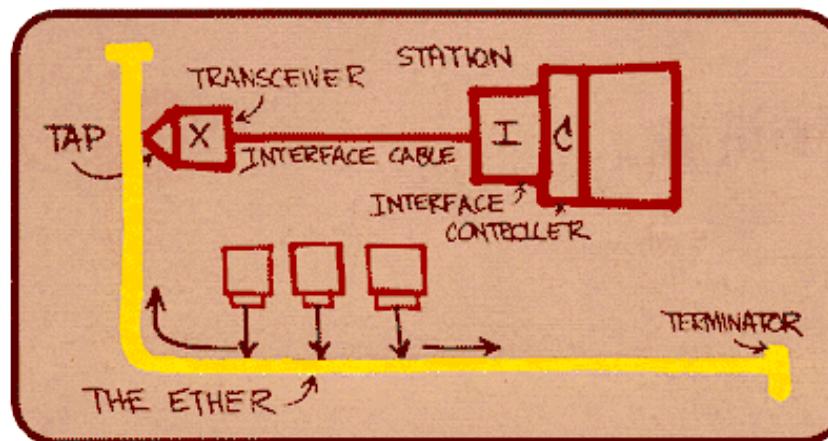
Link layer, LANs: outline

- | | |
|---|---|
| <p>6.1 introduction, services</p> <p>6.2 error detection,
correction</p> <p>6.3 multiple access
protocols</p> <p>6.4 LANs</p> <ul style="list-style-type: none">– addressing, ARP– Ethernet– switches– VLANs | <p>6.5 link virtualization:
MPLS</p> <p>6.6 data center
networking</p> <p>6.7 a day in the life of a
web request</p> |
|---|---|

Ethernet

“dominant” wired LAN technology:

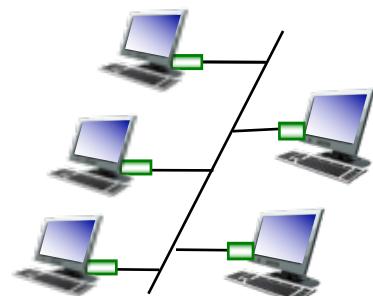
- cheap \$20 for NIC (or less)
- first widely used LAN technology
- simpler, cheaper than token LANs and ATM
- kept up with speed race: 10 Mbps – 100 Gbps



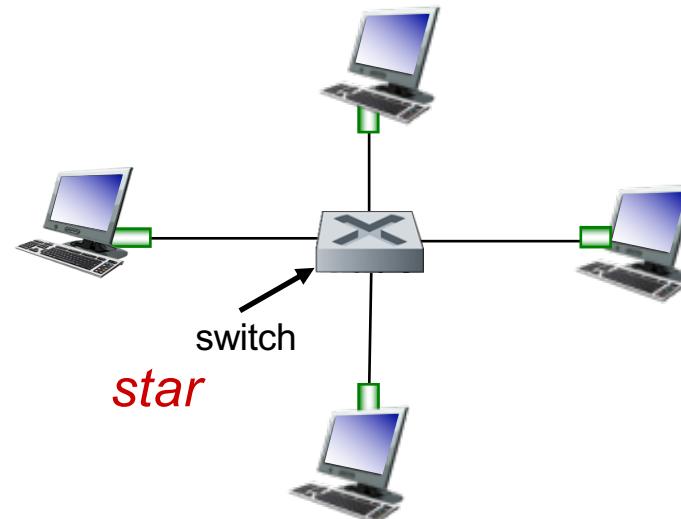
Bob Metcalfe's Ethernet sketch

Ethernet: physical topology

- *bus*: popular through mid 90s
 - all nodes in same collision domain (can collide with each other)
- *star*: prevails today
 - active *switch* in center
 - each “spoke” runs a (separate) Ethernet protocol (nodes do not collide with each other)



bus: coaxial cable





Ethernet frame structure

sending adapter encapsulates IP datagram in **Ethernet frame**



preamble:

- ❖ 7 bytes with pattern 10101010 followed by one byte with pattern 10101011
- ❖ used to synchronize receiver and sender clock rates

Ethernet frame structure (more)

- ❖ **addresses**: 48-bit source and destination MAC addresses
 - if adapter receives frame with matching destination address, or with broadcast address (e.g. ARP packet), it passes data in frame to network layer protocol
 - otherwise, adapter discards frame
- ❖ **type**: indicates higher layer protocol
 - ❖ IPv4, IPv6, ARP, and many more
- ❖ **CRC**: cyclic redundancy check at receiver
 - error detected: frame is dropped



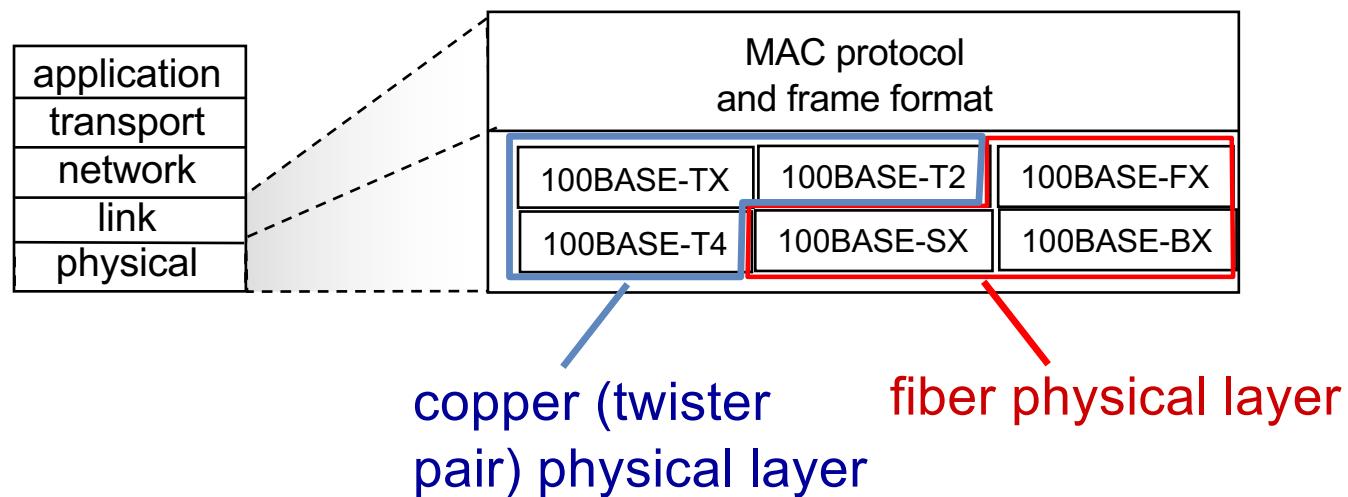


Ethernet: unreliable, connectionless

- *connectionless*: no handshaking between sending and receiving NICs
- *unreliable*: receiving NIC does not send acknowledgements (positive or negative) to sending NIC
- Ethernet's MAC protocol: unslotted *CSMA/CD with binary back-off*
 - Back-off: wait before trying again after a collision
 - > *Time to wait is randomized*
 - > *After c collisions, wait between 0 and $2^c - 1$ time units*

802.3 Ethernet standards: link & physical layers

- ❖ *many* different Ethernet standards
 - common MAC protocol and frame format
 - different speeds: 2 Mbps, 10 Mbps, 100 Mbps, 1Gbps, 10 Gbps, 40 Gbps, 100 Gbps, 200 Gbps, 400 Gbps
 - different physical layer media: fiber, cable





Ethernet Standards (Selected)

Standard	Year	Description
Ethernet II (DIX)	1982	10 Mbit/s over thick coax cable
IEEE 802.3	1983	10 Mbit/s over thick coax cable
IEEE 802.3a	1985	10 Mbit/s over thin coax cable
IEEE 802.3i	1990	10 Mbit/s over twisted pair
IEEE 802.3j	1993	10 Mbit/s over fiber
IEEE 802.3u	1995	100 Mbit/s over twisted pair and fiber (Fast Ethernet)
IEEE 802.3x	1997	Full duplex, flow control, DIX framing
IEEE 802.3z	1998	1 Gbit/s Ethernet over fiber (Gigabit Ethernet)
IEEE 802.3ab	1999	1 Gbit/s Ethernet over twisted pair (Gigabit Ethernet)
IEEE 802.3ae	2003	10 Gbit/s Ethernet over fiber
IEEE 802.3af	2003	Power over Ethernet (12.95 W)
IEEE 802.3ah	2004	Ethernet in the First Mile
IEEE 802.3an	2006	10 Gbit/s Ethernet over twisted pair
IEEE 802.3av	2009	10 Gbit/s EPON
IEEE 802.3az	2010	Energy Efficient Ethernet
IEEE 802.3ba	2010	40 Gbit/s and 100 Gbit/s Ethernet over fiber and (short) copper



Link layer, LANs: outline

- 6.1 introduction, services
- 6.2 error detection,
correction
- 6.3 multiple access
protocols
- 6.4 LANs
 - addressing, ARP
 - Ethernet
 - switches
 - ~~VLANs~~
- 6.5 link virtualization:
MPLS
- 6.6 data center
networking
- 6.7 a day in the life of a
web request

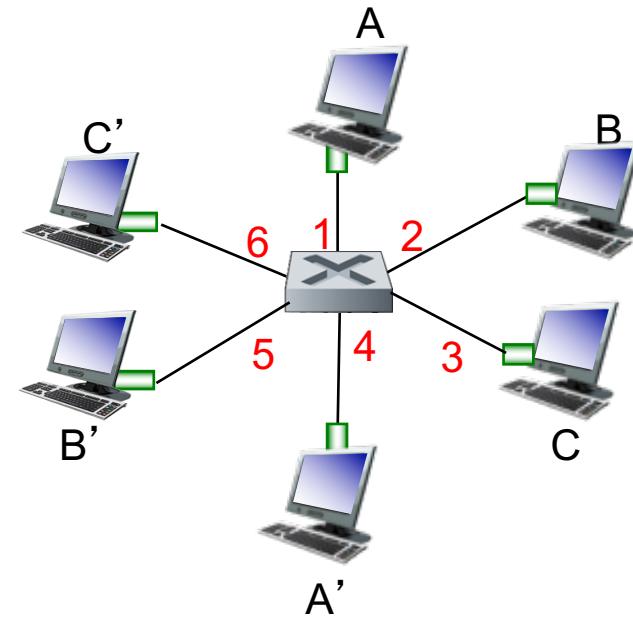


Ethernet switch

- link-layer device: takes an *active* role
 - store and forward Ethernet frames
 - examine incoming frame's MAC address, **selectively** forward frame to one (or more) outgoing links
 - when frame is to be forwarded on segment, uses CSMA/CD to access segment
- *transparent*
 - hosts are unaware of presence of switches
- *plug-and-play, self-learning*
 - switches do not need to be configured

Switch: *multiple simultaneous transmissions*

- each host has dedicated, direct connection to switch
- switches buffer packets
- Ethernet protocol used on each incoming link, but no collisions; *full duplex*
 - each link is its own collision domain
- *switching*: A-to-A' and B-to-B' can transmit simultaneously, without collisions



*switch with six interfaces
(1,2,3,4,5,6)*

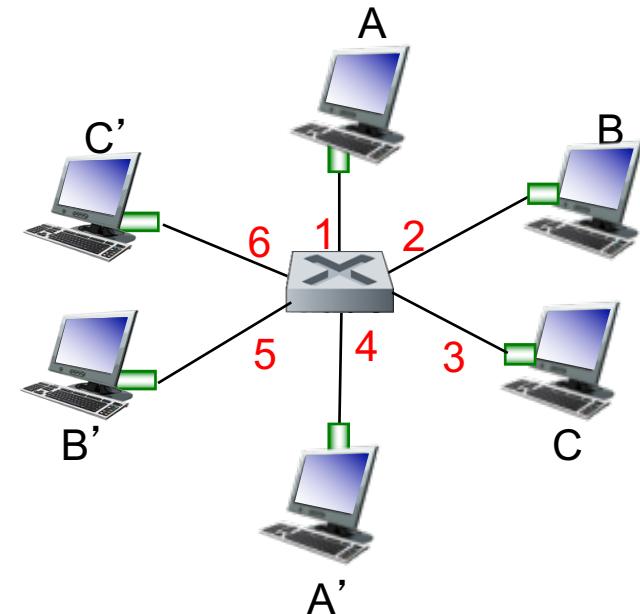
Switch forwarding table

Q: how does switch know that A' is reachable via interface 4, and B' via interface 5?

- ❖ **A:** each switch has a **switch table**, each table entry contains:
 - (MAC address of host, interface to reach host, time stamp)
 - looks like a routing table!

Q: how are entries created and maintained in switch table?

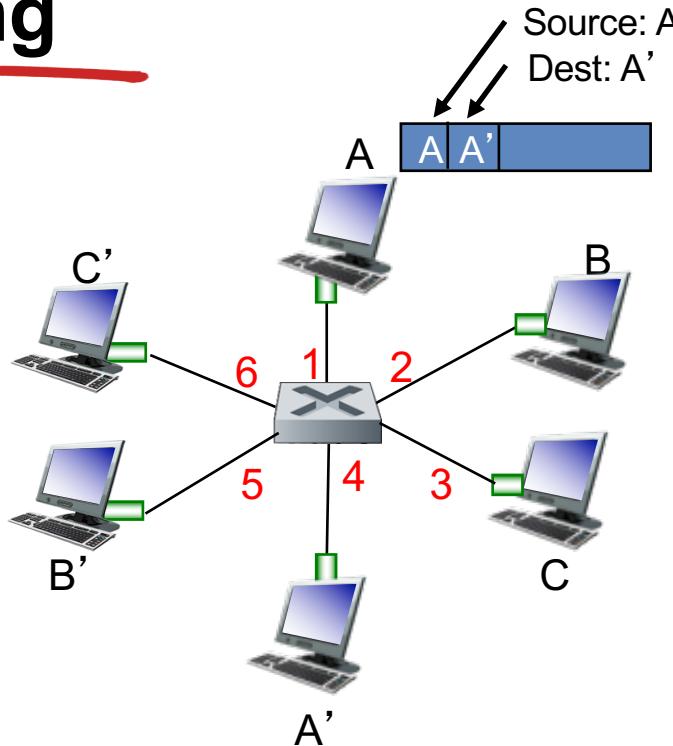
- something like a routing protocol?
- ❖ **A:** No. Through **self-learning**



*switch with six interfaces
(1,2,3,4,5,6)*

Switch: self-learning

- switch *learns* which hosts can be reached through which interfaces
 - when frame received, switch “learns” location of sender: incoming LAN segment
 - records sender/location pair in switch table



*Switch table
(initially empty)*

MAC addr	interface	TTL
A	1	60



Switch: frame filtering/forwarding

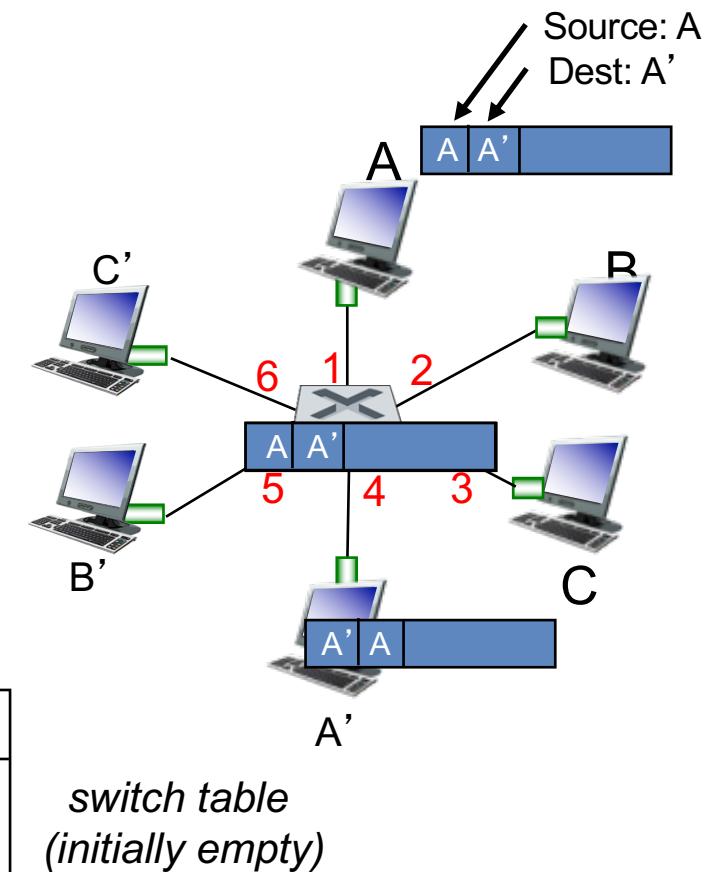
when frame received at switch:

1. record incoming link and MAC address of sending host
2. examine switch table:
3. if destination MAC address found in switch table
then {
 if destination is on segment from which frame arrived
 then drop frame
 else forward frame on interface given by switch table
}
else flood /* forward on all interfaces except arriving
interface */

Self-learning, forwarding: example

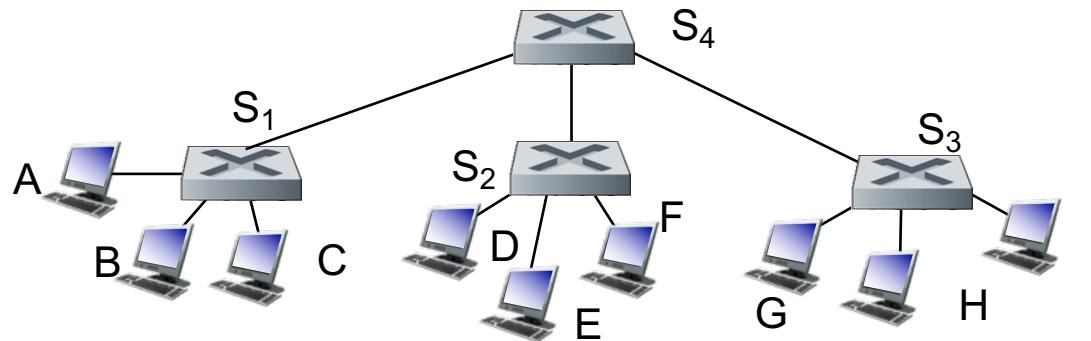
- frame destination, A', location unknown: **flood**
 - ❖ destination A location known: **selectively send on just one link**

MAC addr	interface	TTL
A	1	60
A'	4	60



Interconnecting switches

- ❖ switches can be connected together

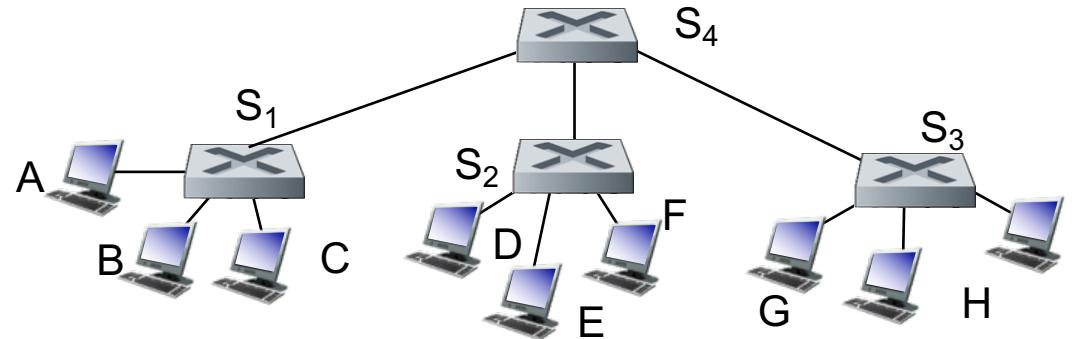


Q: sending from A to G - how does S_1 know to forward frame destined to F via S_4 and S_3 ?

- ❖ A: self learning! (works exactly the same way as in single-switch case!)

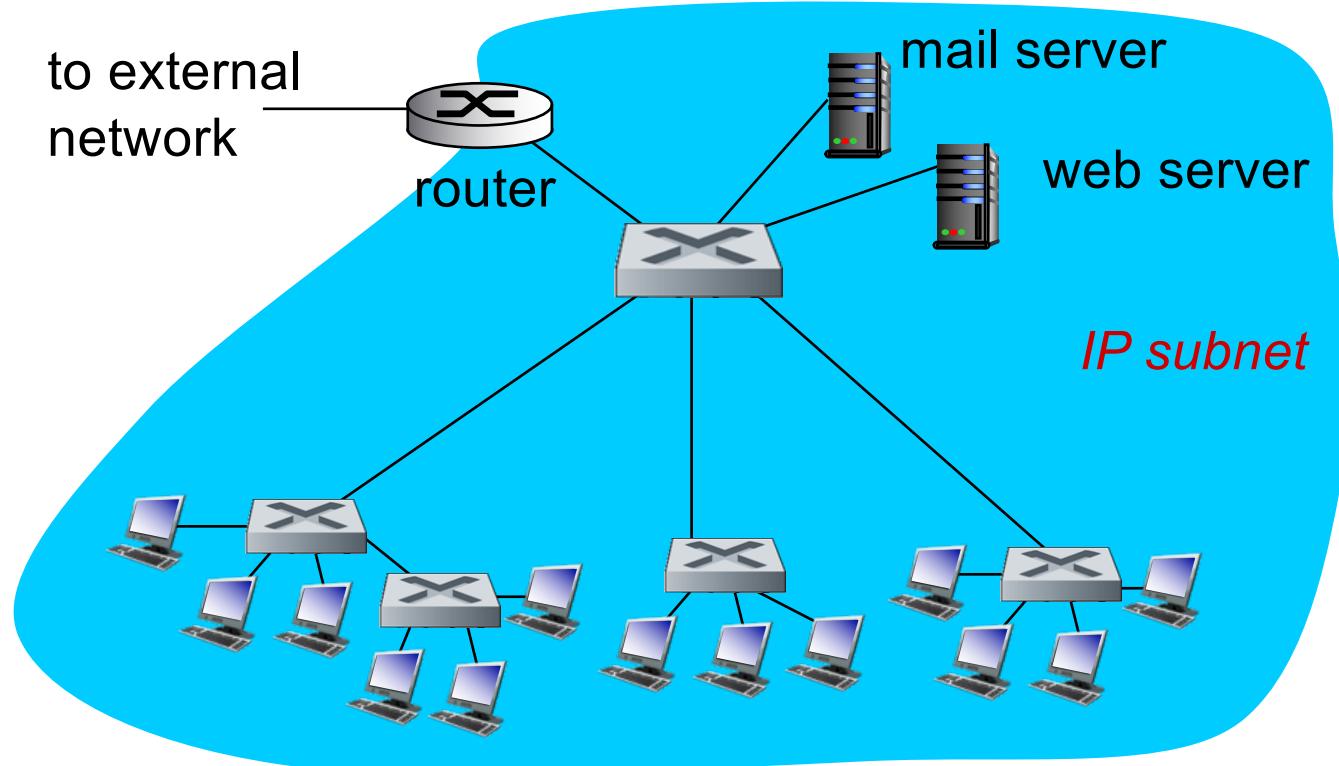
Self-learning multi-switch example

Suppose C sends frame to I, I responds to C



- ❖ **Q:** show switch tables and packet forwarding in S_1, S_2, S_3, S_4

Institutional network



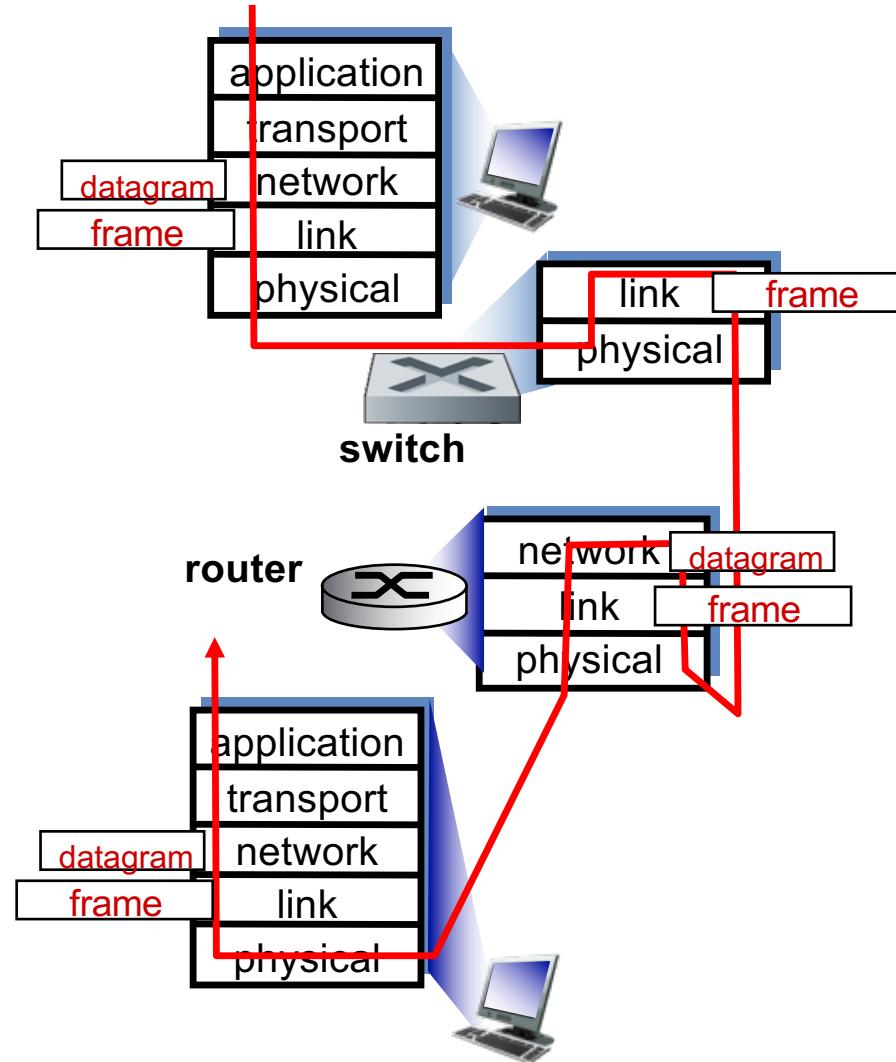
Switches vs. routers

both are store-and-forward:

- **routers**: network-layer devices
(examine network-layer headers)
- **switches**: link-layer devices
(examine link-layer headers)

both have forwarding tables:

- **routers**: compute tables using routing algorithms, IP addresses
- **switches**: learn forwarding table using flooding, learning, MAC addresses





Link layer, LANs: outline

6.1 introduction, services

6.2 error detection,
correction

6.3 multiple access
protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization:

MPLS

6.6 data center
networking

6.7 a day in the life of a
web request



Link layer, LANs: outline

6.1 introduction, services

6.2 error detection,
correction

6.3 multiple access
protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization:
MPLS

6.6 data center
networking

6.7 a day in the life of a
web request

Data center networks

- 10's to 100's of thousands of hosts, often closely coupled, in close proximity:
 - e-business (e.g. Amazon)
 - content-servers (e.g., YouTube, Facebook, Akamai, Apple, Microsoft)
 - search engines, data mining (e.g., Google)
- ❖ challenges:
 - multiple applications, each serving massive numbers of clients
 - managing/balancing load, avoiding bottlenecks in processing, networking, and data access

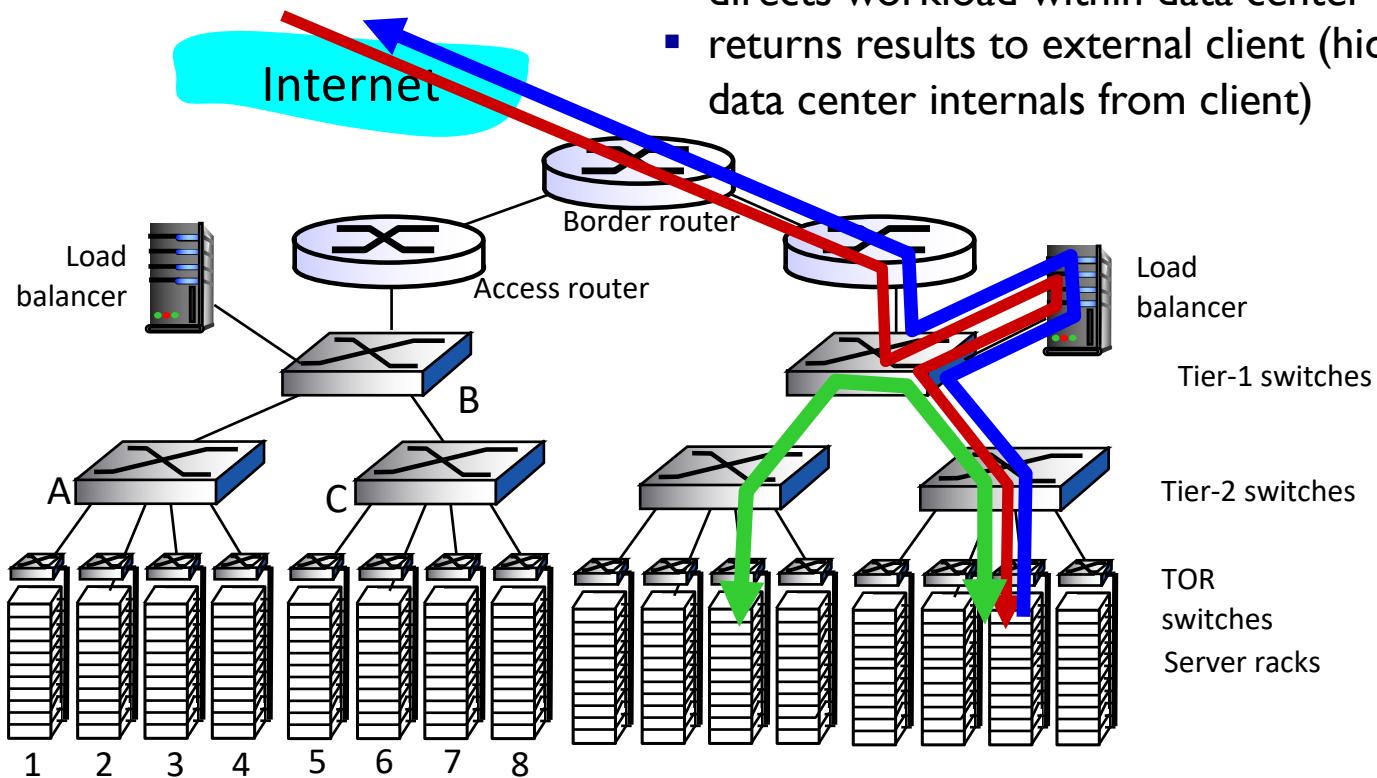


*Inside a 40-ft Microsoft container,
Chicago data center*

Data center networks

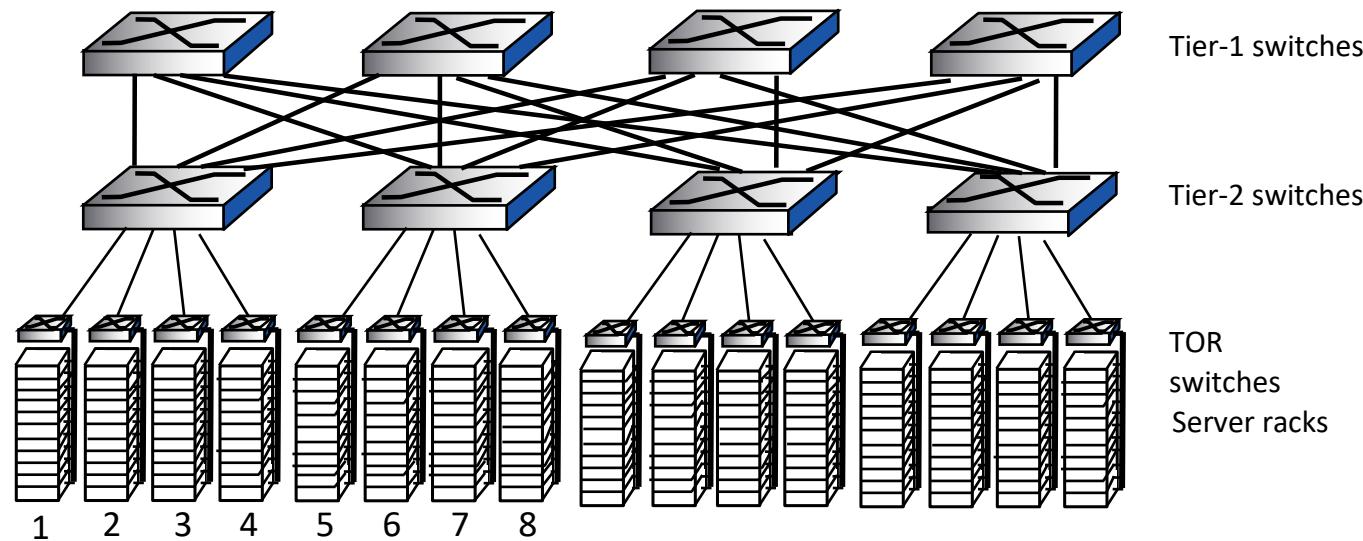
load balancer: application-layer routing

- receives external client requests
- directs workload within data center
- returns results to external client (hiding data center internals from client)



Data center networks

- ❖ rich interconnection among switches, racks:
 - increased throughput between racks (multiple routing paths possible)
 - increased reliability via redundancy





Link layer, LANs: outline

6.1 introduction, services

6.2 error detection,
correction

6.3 multiple access
protocols

6.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

6.5 link virtualization:
MPLS

6.6 data center
networking

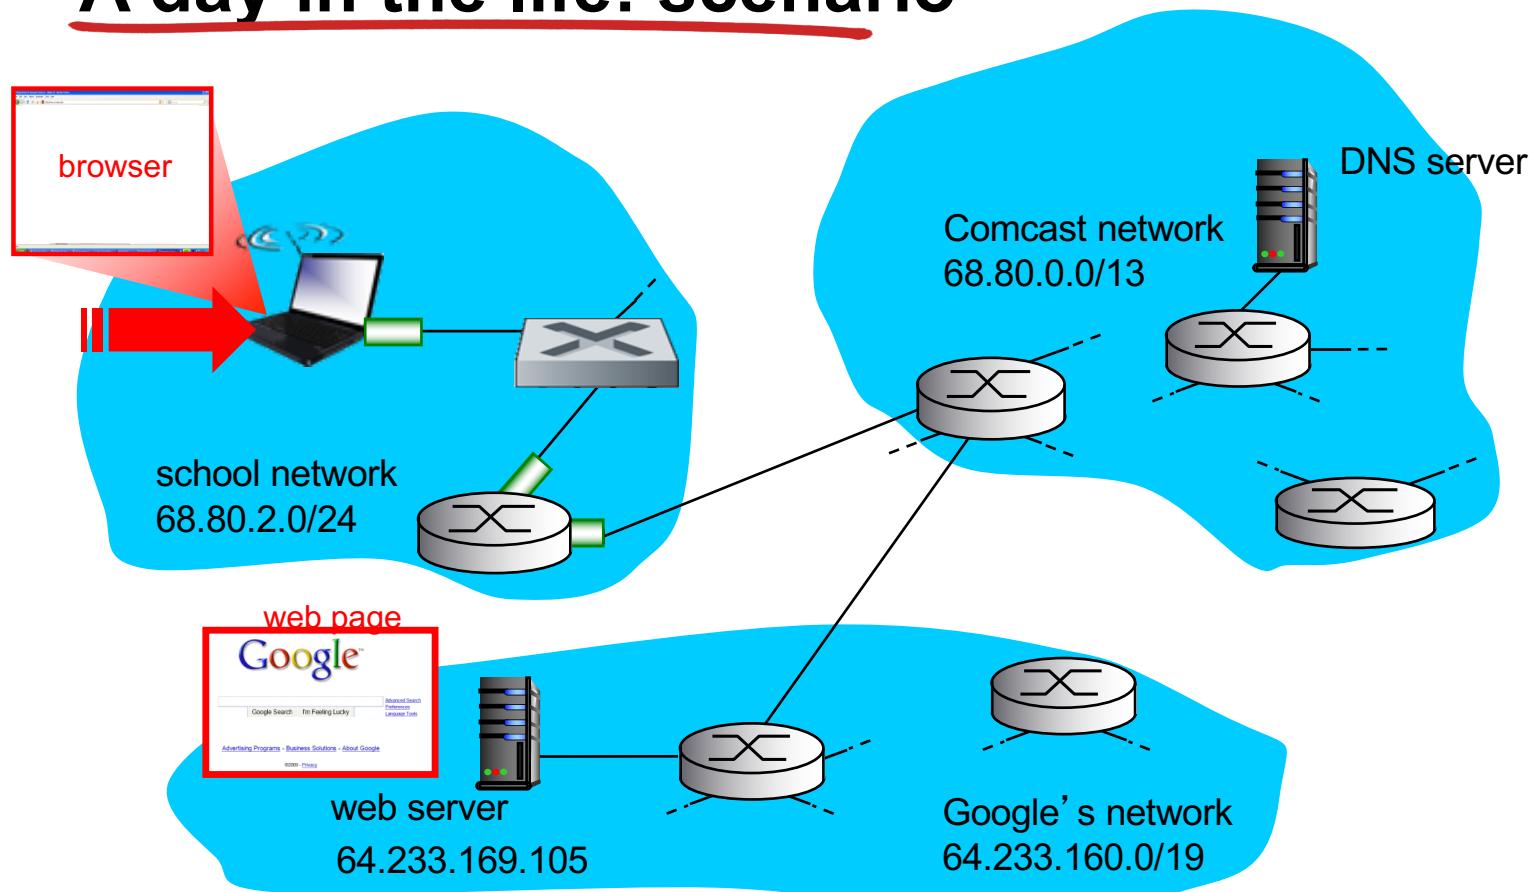
6.7 a day in the life of a
web request



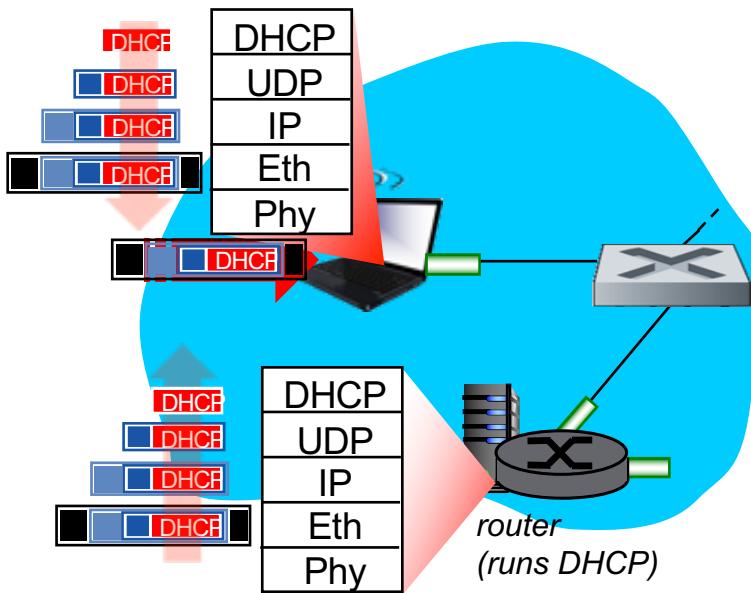
Synthesis: a day in the life of a web request

- ❖ journey down protocol stack complete!
 - application, transport, network, link
- ❖ putting-it-all-together: synthesis!
 - *goal*: identify, review, and understand protocols (at all layers) involved in seemingly simple scenario: requesting www page
 - *scenario*: student attaches laptop to campus network, and requests/receives www.google.com

A day in the life: scenario

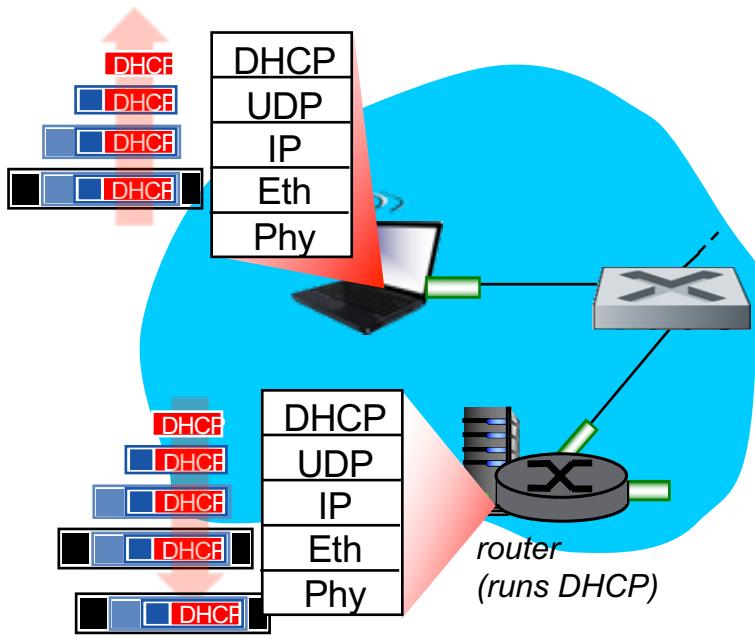


A day in the life... connecting to the Internet



- ❖ connecting laptop needs to get its own IP address, address of first-hop router, and address of DNS server: use **DHCP**
- ❖ DHCP request **encapsulated** in **UDP**, encapsulated in **IP**, encapsulated in **802.3** Ethernet
- ❖ Ethernet frame **broadcast** (destination: FF-FF-FF-FF-FF-FF) on LAN, received at router running **DHCP** server
- ❖ Ethernet **demultiplexed** to IP demultiplexed to UDP demultiplexed to DHCP

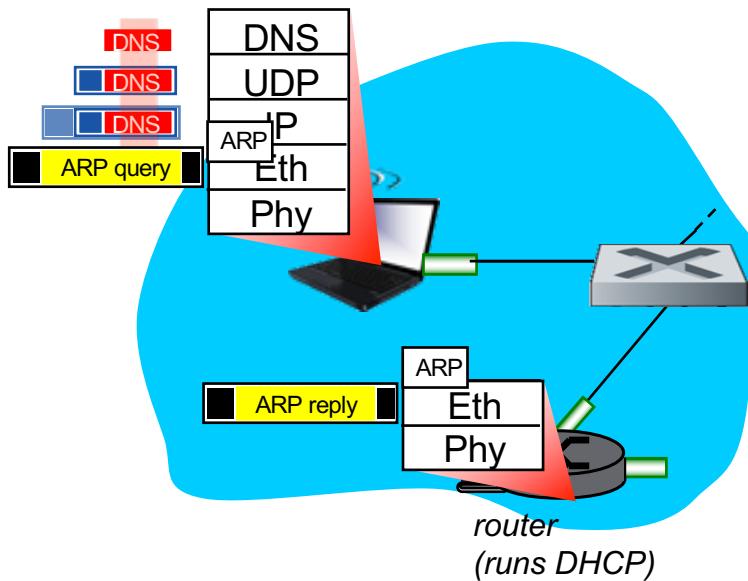
A day in the life... connecting to the Internet



- DHCP server formulates **DHCP ACK** containing client's IP address, IP address of first-hop router for client, and name & IP address of DNS server
- ❖ encapsulation at DHCP server, frame forwarded (**switch learning**) through LAN, demultiplexing at client
- ❖ DHCP client receives DHCP ACK reply

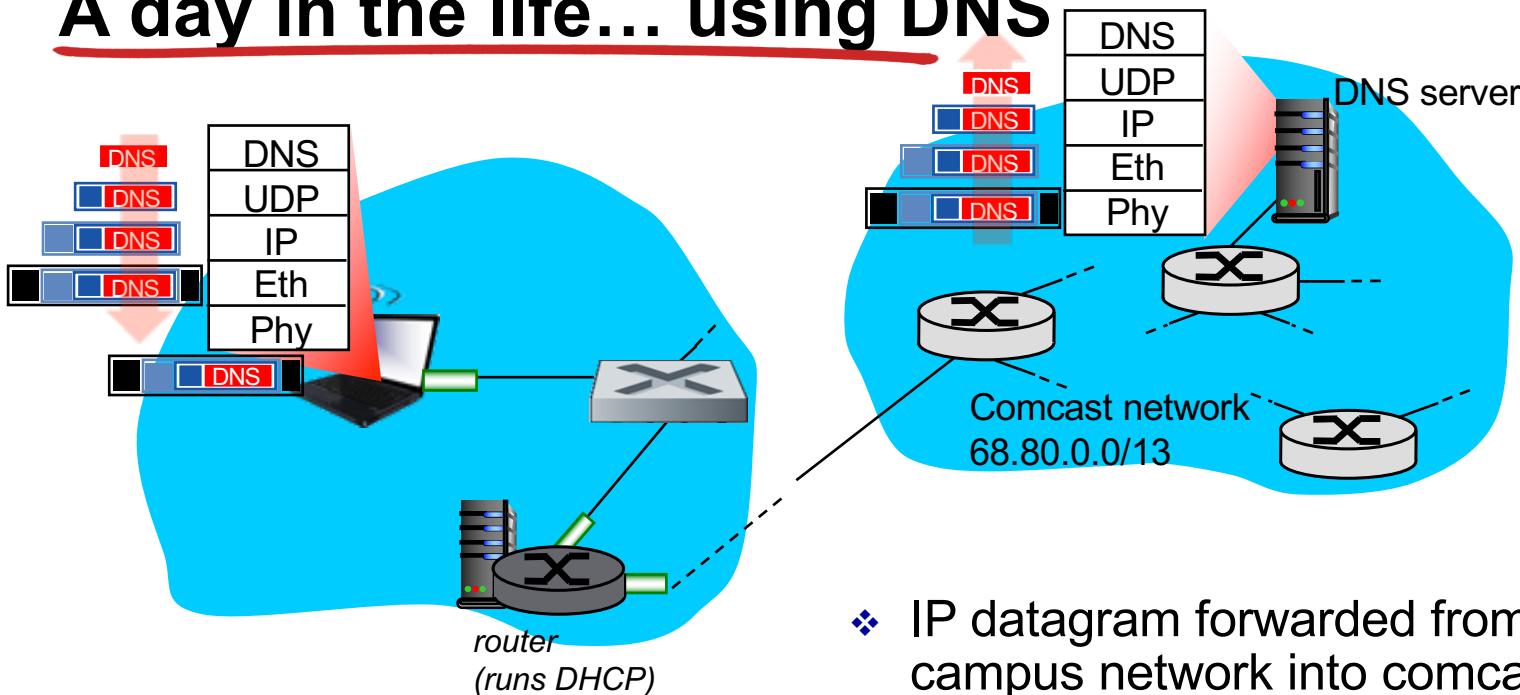
Client now has its IP address, knows name & address of DNS server, and IP address of its first-hop router

A day in the life... ARP (before DNS, before HTTP)



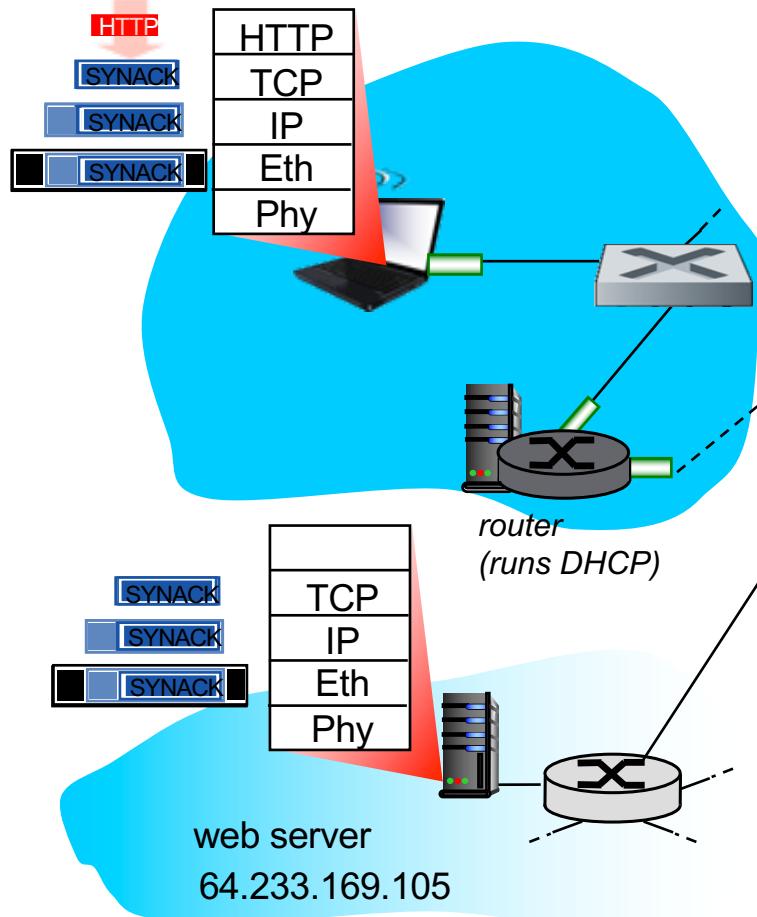
- ❖ before sending *HTTP* request, client needs IP address of www.google.com: **DNS**
- ❖ DNS query created, encapsulated in UDP, encapsulated in IP, encapsulated in Ethernet. To send frame to router, need MAC address of router interface: **ARP**
- ❖ **ARP query** broadcast, received by router, which replies with **ARP reply** giving MAC address of router interface
- ❖ client now knows MAC address of first hop router, so can now send frame containing DNS query

A day in the life... using DNS



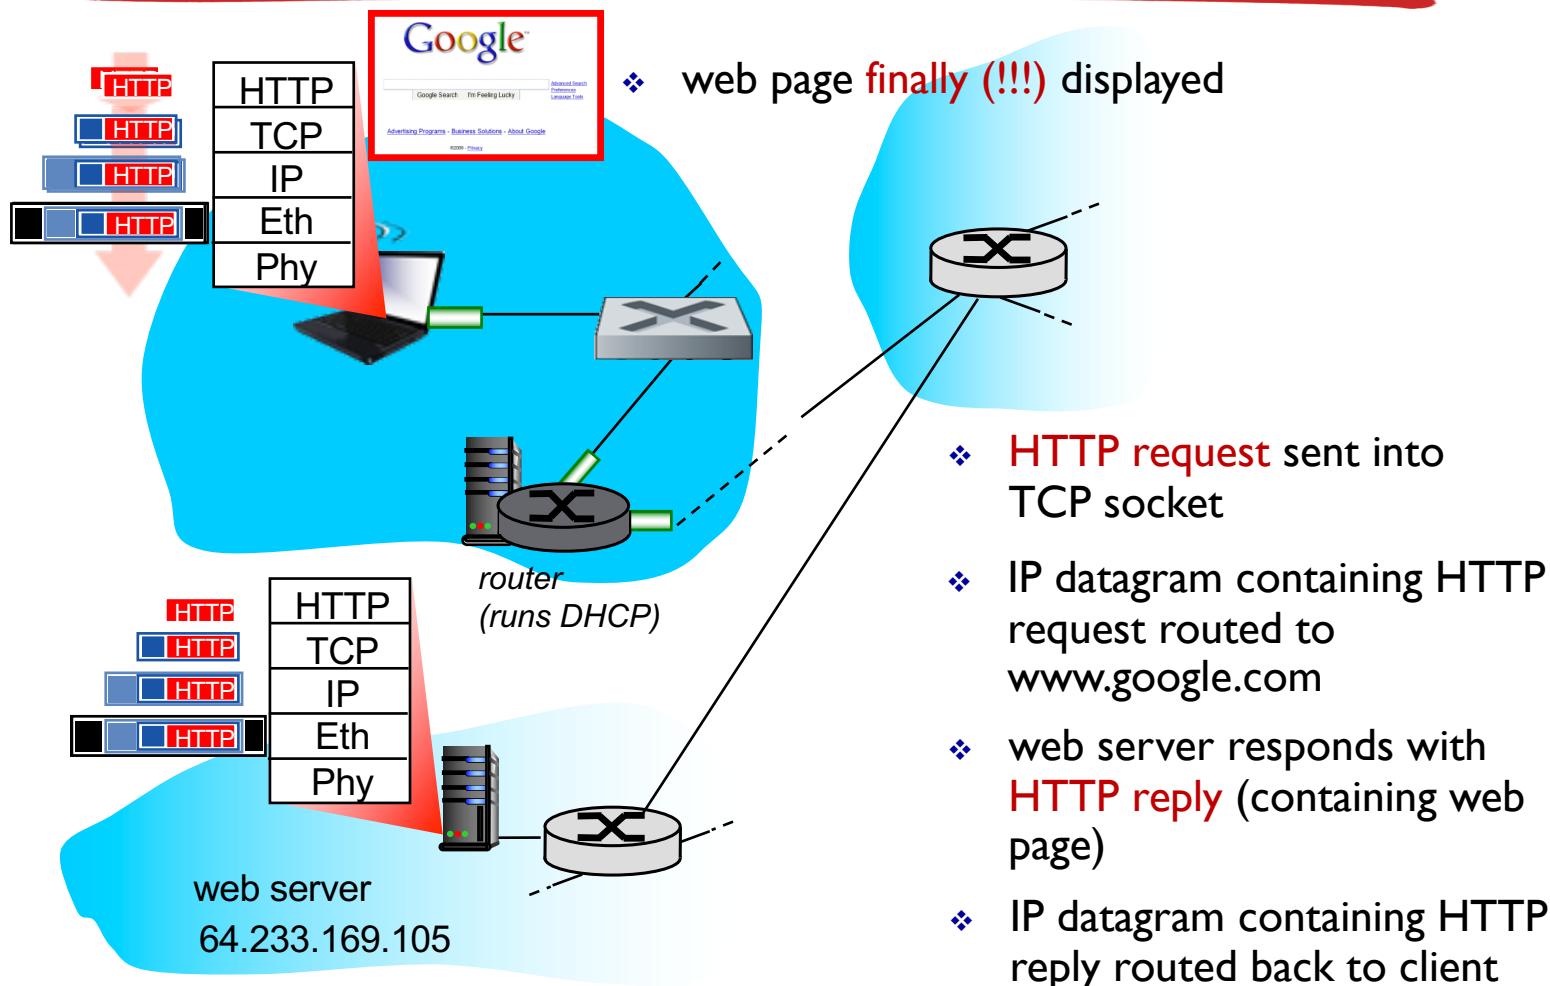
- ❖ IP datagram forwarded from campus network into comcast network, routed (tables created by **RIP, OSPF, IS-IS** and/or **BGP** routing protocols) to DNS server
- ❖ demultiplexed to DNS server
- ❖ DNS server replies to client with IP address of www.google.com

A day in the life...TCP connection carrying HTTP



- ❖ to send HTTP request, client first opens **TCP socket** to web server
- ❖ TCP **SYN segment** (step 1 in 3-way handshake) inter-domain routed to web server
- ❖ web server responds with **TCP SYNACK** (step 2 in 3-way handshake)
- ❖ TCP **connection established!**

A day in the life... HTTP request/reply





Chapter 6: Summary

- ❖ principles behind data link layer services:
 - error detection, correction
 - sharing a broadcast channel: multiple access
 - link layer addressing
- ❖ instantiation and implementation of various link layer technologies
 - Ethernet
 - switched LANS, VLANs
 - virtualized networks as a link layer: MPLS
- ❖ synthesis: a day in the life of a web request



Chapter 6: let's take a breath

- journey down protocol stack *complete* (except PHY)
- solid understanding of networking principles, practice
- could stop here but *lots* of interesting topics!
 - wireless
 - multimedia
 - security
 - network management